NextIgnition — Cost Plan for 1,000-2,000 Users

Server hosting · Frontend hosting · AI/LLM · Email · SMS — sized for 1k-2k users

Prepared June 2026 · USD unless noted · vendor prices change — verify before relying

This plan sizes NextIgnition for a real user base of 1,000-2,000 registered users (assume ~1,500 monthly active, ~300 daily active). It keeps the same lean philosophy as the MVP guide but bumps each service to the tier that actually holds up at this scale. Two end-to-end monthly scenarios are given: LEAN (~1k users, careful) and RECOMMENDED (1k-2k users with headroom). Usage costs (SMS, payments) scale on top and are estimated separately.

1. Workload assumptions at this scale

Metric	Assumption (1k-2k users)	Why it matters
Registered users	1,000 - 2,000	Well within every "medium" tier limit below
Monthly active (MAU)	~1,500	Supabase Pro includes 100,000 MAU — no auth overage
Daily active (DAU)	~250 - 350	Drives AI requests, DB reads, bandwidth
AI requests / day	~1,500 - 3,000	Exceeds the 1,000/day free cap -> some paid overflow
Emails / day (verify + notify)	~150 - 400	Exceeds Resend free 100/day cap -> paid email
SMS OTPs / month	~3,000 - 6,000	Signups + phone logins; pay-per-use
File storage (avatars + decks)	~5 - 30 GB	Comfortably inside Supabase Pro 100 GB

2. AI / LLM costing (per 1,000,000 tokens)

Provider / Model (as used in app)	Input $/1M	Output $/1M	Role at 1k-2k	Notes
OpenRouter "openrouter/free" (DEFAULT)	$0.00	$0.00	Primary	Free 1,000 req/day after one-time $10 credit
Google Gemini 2.5 Flash-Lite	$0.10	$0.40	Paid overflow	Cheapest reliable fallback once daily free cap is hit
Google Gemini 3.1 Flash-Lite	$0.25	$1.50	Optional	GA budget model; better quality
Google Gemini 3.5 Flash	$1.50	$9.00	Premium only	Use only for high-value flows
OpenAI GPT-4 Turbo (useGPT4 = ON)	$10.00	$30.00	Avoid	Most expensive path — keep OFF

Math: ~1.5k-3k AI requests/day. The first 1,000/day stay free (after the one-time $10 OpenRouter credit). The overflow (~500-2,000 req/day) goes to Gemini 2.5 Flash-Lite. At ~2,500 tokens in / ~600 out per request, ~45,000 paid requests/month = ~110M input + ~27M output tokens = ~$11 input + ~$11 output "H $15-30/month. Keep u caching (~90% input savings) to push this toward the low end.

3. Server / backend hosting (per month)

Service	Plan	Cost / mo	Fit for 1k-2k	Notes
Render (backend API)	Starter	$7	Lean / tight	512 MB RAM, always-on — OK for ~1k users, watch memory
Render (backend API)	Standard	$25	Recommended	2 GB RAM — comfortable headroom for 1k-2k + AI calls
Render (backend API)	Pro	$85	Overkill (yet)	4 GB RAM — only if traffic spikes hard
Supabase (DB + storage + auth)	Pro	$25	Recommended	Always-on, backups, 8 GB DB, 100 GB files, 100k MAU
Supabase	Pro + compute	$35+	If DB is hot	Bump compute only if query latency climbs

Service	Plan	Cost / mo	Fit for 1k-2k	Notes
	add-on			

DB lives in Supabase, so no separate Render database is needed. For 1k-2k users the realistic backend is Render Standard ($25) + Supabase Pro ($25). Render Starter ($7) can work for the lean ~1k end but 512 MB RAM is tight once concurrent AI/DB load rises — size up to Standard at the first sign of memory pressure.

4. Frontend hosting (per month)

Service	Free tier	Paid	Commercial on free?	Notes
Cloudflare Pages	$0 (unlimited bandwidth)	~$5-20	Yes	Best pick — free + commercial even at 1k-2k users
Netlify	$0 (100 GB)	$19	Yes (basic)	100 GB likely enough; upgrade if bandwidth grows
Render Static Site	$0	-	Yes	Keeps frontend + backend on one provider
Vercel	$0 (100 GB)	$20 / seat	No	Hobby is non-commercial; a paid app needs Pro ($20)

At 1k-2k users frontend bandwidth is still small. Cloudflare Pages stays $0 and allows commercial use — no reason to pay $20 for Vercel Pro.

5. Email & SMS (volume-driven)

Service	Plan	Cost / mo	Capacity	Notes
Resend (email)	Free	$0	3,000/mo · 100/day	Too tight — 1k-2k users blow past 100/day
Resend (email)	Pro	$20	50,000/mo	Recommended — covers verify + notifications comfortably
Twilio Verify (OTP)	Pay-per-use	~$16-27	3k-6k OTPs	India ~Rs 0.45/OTP incl. forex surcharge
MSG91 / Gupshup (OTP)	Pay-per-use	~$5-9	3k-6k OTPs	India ~Rs 0.15-0.18/OTP — ~60% cheaper than Twilio

Email: move to Resend Pro ($20) — at 1k-2k users daily verification + notification email exceeds the free 100/day cap. SMS: stays pay-per-use. At ~3k-6k OTPs/month Twilio costs ~$16-27; switching India OTP traffic to MSG91/Gupshup cuts that to ~$5-9. Payments (Razorpay) are 2% + GST per successful transaction — a cost of revenue, not a fixed fee.

6. Monthly cost scenarios (1,000-2,000 users)

Component	LEAN (~1k users, careful)	RECOMMENDED (1k-2k, headroom)
Backend (Render)	Starter ($7)	Standard ($25)
Database/Storage/Auth (Supabase)	Pro ($25)	Pro ($25)
Frontend hosting	Cloudflare Pages ($0)	Cloudflare Pages ($0)
AI / LLM (free + paid overflow)	~$10	~$25
Email (Resend)	Free 3k ($0, if <100/day)	Pro 50k ($20)
SMS OTP	MSG91 ~$5-9	Twilio ~$16-27
Payments (Razorpay)	2% + GST / txn	2% + GST / txn
FIXED MONTHLY TOTAL	~ $42 / month	~ $95 / month
ALL-IN incl. usage (est.)	~ $50-60 / month	~ $110-130 / month

LEAN ~ $42/mo fixed: Render Starter ($7) + Supabase Pro ($25) + free frontend + light AI overflow (~$10), keeping email on the free tier by staying under 100/day. Workable near 1,000 users but little headroom. RECOMMENDED ~ $95/mo fixed: Render Standard ($25) + Supabase Pro ($25) + free frontend + AI (~$25) + Resend Pro ($20) — the realistic plan for a live product at